

How Deep Learning and RFdiffusion transformed De novo Protein Engineering

Mai Ha Phuong

Hanoi University of Science and Technology

September 7, 2026

1 Introduction

For over half a century, our understanding of protein folding has been guided by Anfinsen’s thermodynamic hypothesis. In essence, the native folded conformation of a protein corresponds to the global minimum of Gibbs free energy determined by its amino acid sequences. Building directly upon Anfinsen’s hypothesis, the Rosetta software suite emerged as a computational platform for protein structure prediction and design, focused on three primary challenges: backbone sampling, sequence design and energy-based ranking. Despite major achievements, physics-based force field methods encountered severe limitations: Blindness to alternative folded states, rugged and immense conformational landscape, low success rates for unconstrained and complex designs. The limitations of physics-based approaches have driven a paradigm shift toward deep learning. Among these, RFdiffusion-a generative diffusion model fine-tuned from the structure prediction network RoseTTAFold-has emerged as a breakthrough, enabling the design of diverse, large, and experimentally validated proteins. However, despite these advances, significant challenges remain.

2 Problem Statements

2.1 The Anfinsen Hypothesis and the Free Energy Principle

Traditional de novo pretein design rests on Anfinsen’s hypothesis: a protein’s native structure corresponds to the global minimum of Gibbs free energy determined by its amino acid sequence. Based on this thermodynamic principle, the objective of classical computational design is to identify amino acid sequences that allow a target backbone structure to reach its lowest calculated physical energy state. (1)

2.2 Traditional Physics-Based Force Field Approaches-Rosetta

To tackle challenges, conventional design software such as Rosetta executes three core computational steps:

- Backbone Sampling: Sampling the conformational space to generate plausible 3D backbone structures for a given amino acid chain.
- Energy calculation: Utilizing atomic force fields to compute total physical energy
- Optimization and sequence ranking: Searching for low-energy sequence candidates that fit the target backbone and ranking design models based on their total calculated physical energy.

2.3 Limitations

However, this physics-based paradigm faces three fundamental limitations:

- Energy functions are approximations that cannot capture quantum-level interactions with precision
- Computational cost are prohibitive - Rosetta requires nearly 258 seconds per 100-residue design
- The rugged conformational landscape causes optimization to miss alternative folded states that may possess equal or lower energy

Around 2020, the field underwent a major paradigm shift driven by advances in deep learning. Highly accurate structure prediction networks like ALphaFold and RoseTTAFold established a direct mapping from sequence space to 3D structure space.

3 RFdiffusion and ProteinMPNN: Complementary Solutions

The deep learning paradigm for the de novo protein is best exemplified by the complementary pairing of RFdiffusion and ProteinMPNN. While earlier approaches such as Protein Hallucination demonstrated the potential of neural networks for protein design, they were limited to small proteins (<100 residues) and lacked extensive experimental validation. RFdiffusion builds upon these foundations - drawing inspiration from AlphaFold2's recycling mechanism - to enable robust, large-scale structure generation.

3.1 RFdiffusion in backbone generation

3.1.1 Core principle

RFdiffusion is a denoising diffusion probabilistic model fine-tuned from RoseTTAFold. It generates protein backbones by reversing a noise-adding process over 200 timesteps (3), operating on residue frames. Two innovations are critical: fine-tuning from pretrained RoseTTAFold weights, which leverages prior structural knowledge, and self-conditioning-inspired by AlphaFold2’s recycling - where the model conditions on its own previous prediction of the clean structure. The model operates on residue frames and is trained to reconstruct the original structure from noise by minimizing the mean squared error (MSE) between predicted and true frames

3.1.2 Key innovation

First, fine-tuning from pretrained RoseTTAFold weights leverages the network’s prior knowledge of protein structure, dramatically improving performance over training from scratch. Second, self-conditioning - inspired by AlphaFold2’s recycling mechanism - allows the model to condition on its own previous prediction of the clean structure at each denoising step, enhancing coherence and reducing accumulated errors.(3) Crucially, MSE is used instead of FAPE because it preserves a consistent global reference frame across diffusion steps, which is essential for unconditional generation.

3.1.3 Strength

RFdiffusion generates protein up to 600 residues with high accuracy.(3) It produces novel folds absent from the PDB and supports conditioning on symmetry, binding targets, or functional motifs.

3.1.4 Limitations

RFdiffusion cannot explicitly model small molecules or ligands, which limits enzyme active site design. It also requires substantial computational resources.

3.2 ProteinMPNN: Sequence Design

3.2.1 Core Principle

ProteinMPNN solves the inverse problem: given a protein backbone, predict amino acid sequences that will fold into that structure. It employs a Message Passing Neural Network (MPNN) that represents the protein as a graph - nodes corresponds to residues, and edges encode spatial relationships derived from backbone geometry.(2) (5)

3.2.2 Key innovation

Message passing neural network: Through iterative message passing, each residue aggregates information from its spatial neighbors, building a context-aware representation. The decoder then predicts amino acid identities in an autoregressive manner, allowing sequences to be generated in any order.(2)

3.2.3 Strengths

ProteinMPNN achieves a 52.4% sequence recovery rate on native backbones - dramatically outperforming Rosetta’s 32.9 %. It is over 200 times faster. It requires no expert knowledge to operate and can rescue previously failed designs made using Rosetta or AlphaFold. (2)

3.2.4 Limitations

ProteinMPNN is fundamentally dependent on the quality of the input backbone. It also does not model protein dynamics or conformational changes, and its prediction may exhibit overconfidence (2)

4 Discussion

Despite the transformative success of the RFDiffusion and ProteinMPNN, critical limitations remain that constrain the scope of the de novo protein design. RFDiffusion cannot explicitly model small molecules or ligands during diffusion; enzyme active site design currently relies on manually designed external potentials (combines repulsive and attractive fields) to create complementary pockets rather than co-designing protein and ligands atoms jointly. The model also focuses strictly on protein backbones, leaving nucleic acid - binding proteins and folded RNA structures outside its scope.(3) Its binder design campaigns primarily targeted non-polar interfaces, with lower success rates on polar targets or sites lacking native binding partners. ProteinMPNN’s limitations are intimately tied to its dependence on the input backbone. It cannot account for varying backbone quality at a per-position level. In silico metrics do not guarantee experimental success: a single mutation can abolish folding despite high overall recovery. The model also lacks physical transparency, operating probabilistically rather than optimizing sidechain rotamers or energy minima as Rosetta does.(1) Critically, it operates on a fixed backbone and cannot optimize backbone flexibility during sequence generation. RoseTTAFold All-Atom (RFAA) address these gaps by combining residue-based representations of amino acids and DNA bases with atomic representation of all other groups, enabling modeling of proteins, nucleic acids, small molecules and covalent modifications. Fine-tuned on denoising tasks, RFAA builds protein structures around small molecules by directly

generating binding pockets at atomic resolution, enabling co-design of protein and ligand. It has been experimentally validated for targets including digoxigenin and heme.⁽⁴⁾ More recent iterations, such as RFDiffusion3, extend this paradigm to nucleic-acid binders and complex enzymes active sites, representing a promising trajectory toward closing the structure-function gap

References

- [1] Yang W, Wang S, Lee GR, Zhang JZ, Courbet A, Juergens D, et al. The past, present and future of de novo protein design. *Nature*. 2026;652(8112). doi:10.1038/s41586-026-10328-7.
- [2] Dauparas J, Anishchenko I, Bennett N, Bai H, Ragotte RJ, Milles LF, et al. Robust deep learning-based protein sequence design using ProteinMPNN. *Science*. 2022;378(6615):49–56. doi:10.1126/science.add2187.
- [3] Watson JL, Juergens D, Bennett NR, Trippe BL, Yim J, Eisenach HE, et al. De novo design of protein structure and function with RFDiffusion. *Nature*. 2023;620(7976):1089–1100. doi:10.1038/s41586-023-06415-8.
- [4] Krishna R, Wang J, Ahern W, Sturmfels P, Venkatesh P, Kalvet I, et al. Generalized biomolecular modeling and design with RoseTTAFold All-Atom. *Science*. 2024;384(6693):eadl2528. doi:10.1126/science.adl2528.
- [5] Ingraham J, Garg VK, Barzilay R, Jaakkola T. Generative models for graph-based protein design. In: Wallach H, Larochelle H, Beygelzimer A, d’Alché-Buc F, Fox E, Garnett R, editors. *Advances in Neural Information Processing Systems 32 (NeurIPS 2019)*. Neural Information Processing Systems Foundation; 2019. p. 15741–15752.